

Evaluation

May 27, 2020

1 Evaluation of BERT and DistilBERT

1.1 METEOR and GLEU scores

1. Functions for calculating GLEU and METEOR for two outputs at the same time
2. Calculating GLEU scores
3. Calculating METEOR scores
4. Appending to dataframe, getting 50 samples

```
[7]: import nltk
import json
import re
import numpy
import pandas as pd
import string

from nltk.translate import gleu_score
from nltk.stem.porter import PorterStemmer
from nltk.corpus import wordnet
from itertools import chain, product
```

1.2 1. Functions for calculating GLEU and METEOR for two outputs at the same time

The original GLEU function compares one translation to several golden standards, but we want to compare the performance of both BERT and DistilBERT at the same time to the golden standard.

GLEU takes an input of a split string. METEOR takes just strings.

1.2.1 1.1 GLEU functions

```
[8]: # Function that expands the gleu sentence calculations to be done on more than 1
    ↪ 1 input texts
def gleu_lists_no_df(golden_standard, text1, text2):

    # n-gram lists for text1 and text2
    ngram_list = [[], []]

    # Append the calculated gleu scores to a list for both text1 and text2
```

```

    ngram_list[0].append(nltk.translate.gleu_score.
→sentence_gleu([golden_standard], text1))
    ngram_list[1].append(nltk.translate.gleu_score.
→sentence_gleu([golden_standard], text2))

    # This one returns ngram_list directly for the function that makes a row
→per answer
    return(ngram_list)

# Function that takes columns as input and outputs a dataframe of the results
→per model
def gleu_more(gold1, gold2, berts, distilberts):
    # Initiate results dataframe
    results = pd.DataFrame()

    # Loop over the rows in the dataset
    for i in range(0, len(gold1)):
        # Define the inputs to the bleu function from the dataset
        # One row per input
        gold_1 = gold1[i]
        # Clean of punctuation
        gold_1 = gold_1.translate(str.maketrans('', '', string.punctuation))
        gold_2 = gold2[i]
        #gold_2 = gold_2.translate(str.maketrans('', '', string.punctuation))
        bert = berts[i]
        bert = bert.translate(str.maketrans('', '', string.punctuation))
        distil = distilberts[i]
        distil = distil.translate(str.maketrans('', '', string.punctuation))

        # Get the gleu scores per line in dataframe with the predefined
→function for 1st golden answers
        ngram_list1 = gleu_lists_no_df(gold_1.split(), bert.split(), distil.
→split())

        # Now get the glue scores for 2nd golden answer in case there are more
→than 1 ways to answer the question
        # For missing values in Gold2, return all 0-s
        if pd.isnull(gold_2):
            ngram_list2 = [[0], [0]]
        else:
            ngram_list2 = gleu_lists_no_df(gold_2.split(), bert.split(), distil.
→split())

        # Get the highest value per model in terms of them matching best to
→gold1 or gold2
        ngram_max = [max(i, j) for i, j in zip(ngram_list1, ngram_list2)]

```

```

# Append to dataframe, index is model name + iteration
dff = pd.DataFrame(ngram_max, index=['0', '1'])

# Append results to the dataframe
results = results.append(dff)

return(results)

```

1.2.2 1.2 METEOR functions

We are using other parameter values for meteor score calculations so I will add the whole code on compiling meteor scores together with how the parameter values were changed. Meteor score code retrieved from: https://www.nltk.org/_modules/nltk/translate/meteor_score.html

```

[20]: from nltk.stem.porter import PorterStemmer
from nltk.corpus import wordnet
from itertools import chain, product

def _generate_enums(hypothesis, reference, preprocess=str.lower):
    """
    Takes in string inputs for hypothesis and reference and returns
    enumerated word lists for each of them

    :param hypothesis: hypothesis string
    :type hypothesis: str
    :param reference: reference string
    :type reference: str
    :preprocess: preprocessing method (default str.lower)
    :type preprocess: method
    :return: enumerated words list
    :rtype: list of 2D tuples, list of 2D tuples
    """
    hypothesis_list = list(enumerate(preprocess(hypothesis).split()))
    reference_list = list(enumerate(preprocess(reference).split()))
    return hypothesis_list, reference_list

def exact_match(hypothesis, reference):
    """
    matches exact words in hypothesis and reference
    and returns a word mapping based on the enumerated
    word id between hypothesis and reference

    :param hypothesis: hypothesis string
    :type hypothesis: str

```

```

:param reference: reference string
:type reference: str
:return: enumerated matched tuples, enumerated unmatched hypothesis tuples,
        enumerated unmatched reference tuples
:rtype: list of 2D tuples, list of 2D tuples, list of 2D tuples
"""
hypothesis_list, reference_list = _generate_enums(hypothesis, reference)
return _match_enums(hypothesis_list, reference_list)

def _match_enums(enum_hypothesis_list, enum_reference_list):
    """
    matches exact words in hypothesis and reference and returns
    a word mapping between enum_hypothesis_list and enum_reference_list
    based on the enumerated word id.

    :param enum_hypothesis_list: enumerated hypothesis list
    :type enum_hypothesis_list: list of tuples
    :param enum_reference_list: enumerated reference list
    :type enum_reference_list: list of 2D tuples
    :return: enumerated matched tuples, enumerated unmatched hypothesis tuples,
            enumerated unmatched reference tuples
    :rtype: list of 2D tuples, list of 2D tuples, list of 2D tuples
    """
    word_match = []
    for i in range(len(enum_hypothesis_list))[:-1]:
        for j in range(len(enum_reference_list))[:-1]:
            if enum_hypothesis_list[i][1] == enum_reference_list[j][1]:
                word_match.append(
                    (enum_hypothesis_list[i][0], enum_reference_list[j][0])
                )
                (enum_hypothesis_list.pop(i)[1], enum_reference_list.pop(j)[1])
                break
    return word_match, enum_hypothesis_list, enum_reference_list

def _enum_stem_match(
    enum_hypothesis_list, enum_reference_list, stemmer=PorterStemmer()
):
    """
    Stems each word and matches them in hypothesis and reference
    and returns a word mapping between enum_hypothesis_list and
    enum_reference_list based on the enumerated word id. The function also
    returns a enumerated list of unmatched words for hypothesis and reference.

    :param enum_hypothesis_list:

```

```

: type enum_hypothesis_list:
: param enum_reference_list:
: type enum_reference_list:
: param stemmer: nltk.stem.api.StemmerI object (default PorterStemmer())
: type stemmer: nltk.stem.api.StemmerI or any class that implements a stem_
↪method
: return: enumerated matched tuples, enumerated unmatched hypothesis tuples,
        enumerated unmatched reference tuples
:rtype: list of 2D tuples, list of 2D tuples, list of 2D tuples
"""
stemmed_enum_list1 = [
    (word_pair[0], stemmer.stem(word_pair[1])) for word_pair in_
↪enum_hypothesis_list
]

stemmed_enum_list2 = [
    (word_pair[0], stemmer.stem(word_pair[1])) for word_pair in_
↪enum_reference_list
]

word_match, enum_unmat_hypo_list, enum_unmat_ref_list = _match_enums(
    stemmed_enum_list1, stemmed_enum_list2
)

enum_unmat_hypo_list = (
    list(zip(*enum_unmat_hypo_list)) if len(enum_unmat_hypo_list) > 0 else_
↪[]
)

enum_unmat_ref_list = (
    list(zip(*enum_unmat_ref_list)) if len(enum_unmat_ref_list) > 0 else []
)

enum_hypothesis_list = list(
    filter(lambda x: x[0] not in enum_unmat_hypo_list, enum_hypothesis_list)
)

enum_reference_list = list(
    filter(lambda x: x[0] not in enum_unmat_ref_list, enum_reference_list)
)

return word_match, enum_hypothesis_list, enum_reference_list

def stem_match(hypothesis, reference, stemmer=PorterStemmer()):
    """
    Stems each word and matches them in hypothesis and reference

```

and returns a word mapping between hypothesis and reference

```
:param hypothesis:
:type hypothesis:
:param reference:
:type reference:
:param stemmer: nltk.stem.api.StemmerI object (default PorterStemmer())
:type stemmer: nltk.stem.api.StemmerI or any class that
implements a stem method
:return: enumerated matched tuples, enumerated unmatched hypothesis tuples,
enumerated unmatched reference tuples
:rtype: list of 2D tuples, list of 2D tuples, list of 2D tuples
"""

enum_hypothesis_list, enum_reference_list = _generate_enums(hypothesis,
↳reference)
    return _enum_stem_match(enum_hypothesis_list, enum_reference_list,
↳stemmer=stemmer)

def _enum_wordnetsyn_match(enum_hypothesis_list, enum_reference_list,
↳wordnet=wordnet):
    """
Matches each word in reference to a word in hypothesis
if any synonym of a hypothesis word is the exact match
to the reference word.

:param enum_hypothesis_list: enumerated hypothesis list
:param enum_reference_list: enumerated reference list
:param wordnet: a wordnet corpus reader object (default nltk.corpus.wordnet)
:type wordnet: WordNetCorpusReader
:return: list of matched tuples, unmatched hypothesis list, unmatched
↳reference list
:rtype: list of tuples, list of tuples, list of tuples

"""
word_match = []
for i in range(len(enum_hypothesis_list))[:-1]:
    hypothesis_syns = set(
        chain(
            *[
                [
                    lemma.name()
                    for lemma in synset.lemmas()
                    if lemma.name().find("_") < 0
                ]
                for synset in wordnet.synsets(enum_hypothesis_list[i][1])
            ]
        )
    )
```

```

    ]
    )
    ).union({enum_hypothesis_list[i][1]})
    for j in range(len(enum_reference_list))[:-1]:
        if enum_reference_list[j][1] in hypothesis_syns:
            word_match.append(
                (enum_hypothesis_list[i][0], enum_reference_list[j][0])
            )
            enum_hypothesis_list.pop(i), enum_reference_list.pop(j)
            break
    return word_match, enum_hypothesis_list, enum_reference_list

def wordnetsyn_match(hypothesis, reference, wordnet=wordnet):
    """
    Matches each word in reference to a word in hypothesis if any synonym
    of a hypothesis word is the exact match to the reference word.

    :param hypothesis: hypothesis string
    :param reference: reference string
    :param wordnet: a wordnet corpus reader object (default nltk.corpus.wordnet)
    :type wordnet: WordNetCorpusReader
    :return: list of mapped tuples
    :rtype: list of tuples
    """
    enum_hypothesis_list, enum_reference_list = _generate_enums(hypothesis,
↳reference)
    return _enum_wordnetsyn_match(
        enum_hypothesis_list, enum_reference_list, wordnet=wordnet
    )

def _enum_allign_words(
    enum_hypothesis_list, enum_reference_list, stemmer=PorterStemmer(),
↳wordnet=wordnet
):
    """
    Aligns/matches words in the hypothesis to reference by sequentially
    applying exact match, stemmed match and wordnet based synonym match.
    in case there are multiple matches the match which has the least number
    of crossing is chosen. Takes enumerated list as input instead of
    string input

    :param enum_hypothesis_list: enumerated hypothesis list
    :param enum_reference_list: enumerated reference list
    :param stemmer: nltk.stem.api.StemmerI object (default PorterStemmer())

```

```

        :type stemmer: nltk.stem.api.StemmerI or any class that implements a stem_
↪method
        :param wordnet: a wordnet corpus reader object (default nltk.corpus.wordnet)
        :type wordnet: WordNetCorpusReader
        :return: sorted list of matched tuples, unmatched hypothesis list,
                unmatched reference list
        :rtype: list of tuples, list of tuples, list of tuples
        """
        exact_matches, enum_hypothesis_list, enum_reference_list = _match_enums(
            enum_hypothesis_list, enum_reference_list
        )

        stem_matches, enum_hypothesis_list, enum_reference_list = _enum_stem_match(
            enum_hypothesis_list, enum_reference_list, stemmer=stemmer
        )

        wns_matches, enum_hypothesis_list, enum_reference_list =
↪_enum_wordnetsyn_match(
            enum_hypothesis_list, enum_reference_list, wordnet=wordnet
        )

        return (
            sorted(
                exact_matches + stem_matches + wns_matches, key=lambda wordpair:
↪wordpair[0]
            ),
            enum_hypothesis_list,
            enum_reference_list,
        )

def align_words(hypothesis, reference, stemmer=PorterStemmer(),
↪wordnet=wordnet):
    """
    Aligns/matches words in the hypothesis to reference by sequentially
    applying exact match, stemmed match and wordnet based synonym match.
    In case there are multiple matches the match which has the least number
    of crossing is chosen.

    :param hypothesis: hypothesis string
    :param reference: reference string
    :param stemmer: nltk.stem.api.StemmerI object (default PorterStemmer())
    :type stemmer: nltk.stem.api.StemmerI or any class that implements a stem_
↪method
    :param wordnet: a wordnet corpus reader object (default nltk.corpus.wordnet)
    :type wordnet: WordNetCorpusReader

```



```

        :return: sorted list of matched tuples, unmatched hypothesis list,
        ↪ unmatched reference list
        :rtype: list of tuples, list of tuples, list of tuples
        """
        enum_hypothesis_list, enum_reference_list = _generate_enums(hypothesis,
        ↪ reference)
        return _enum_align_words(
            enum_hypothesis_list, enum_reference_list, stemmer=stemmer,
        ↪ wordnet=wordnet
        )

def _count_chunks(matches):
    """
    Counts the fewest possible number of chunks such that matched unigrams
    of each chunk are adjacent to each other. This is used to calculate the
    fragmentation part of the metric.

    :param matches: list containing a mapping of matched words (output of
    ↪ _align_words)
    :return: Number of chunks a sentence is divided into post alignment
    :rtype: int
    """
    i = 0
    chunks = 1
    while i < len(matches) - 1:
        if (matches[i + 1][0] == matches[i][0] + 1) and (
            matches[i + 1][1] == matches[i][1] + 1
        ):
            i += 1
            continue
        i += 1
        chunks += 1
    return chunks

def single_meteor_score(
    reference,
    hypothesis,
    preprocess=str.lower,
    stemmer=PorterStemmer(),
    wordnet=wordnet,
    alpha=0.9,
    beta=3,
    gamma=0.5,
):

```

```

"""
Calculates METEOR score for single hypothesis and reference as per
"Meteor: An Automatic Metric for MT Evaluation with High Levels of
Correlation with Human Judgments" by Alon Lavie and Abhaya Agarwal,
in Proceedings of ACL.
http://www.cs.cmu.edu/~alavie/METEOR/pdf/Lavie-Agarwal-2007-METEOR.pdf

>>> hypothesis1 = 'It is a guide to action which ensures that the military_
↳always obeys the commands of the party'

>>> reference1 = 'It is a guide to action that ensures that the military_
↳will forever heed Party commands'

>>> round(single_meteor_score(reference1, hypothesis1),4)
0.7398

    If there is no words match during the alignment the method returns the
    score as 0. We can safely return a zero instead of raising a
    division by zero error as no match usually implies a bad translation.

>>> round(meteor_score('this is a cat', 'non matching hypothesis'),4)
0.0

:param references: reference sentences
:type references: list(str)
:param hypothesis: a hypothesis sentence
:type hypothesis: str
:param preprocess: preprocessing function (default str.lower)
:type preprocess: method
:param stemmer: nltk.stem.api.StemmerI object (default PorterStemmer())
:type stemmer: nltk.stem.api.StemmerI or any class that implements a stem_
↳method
:param wordnet: a wordnet corpus reader object (default nltk.corpus.wordnet)
:type wordnet: WordNetCorpusReader
:param alpha: parameter for controlling relative weights of precision and_
↳recall.
:type alpha: float
:param beta: parameter for controlling shape of penalty as a
               function of as a function of fragmentation.
:type beta: float
:param gamma: relative weight assigned to fragmentation penalty.
:type gamma: float
:return: The sentence-level METEOR score.
:rtype: float
"""

```

```

enum_hypothesis, enum_reference = _generate_enums(
    hypothesis, reference, preprocess=preprocess
)
translation_length = len(enum_hypothesis)
reference_length = len(enum_reference)
matches, _, _ = _enum_align_words(enum_hypothesis, enum_reference,
↪stemmer=stemmer)
matches_count = len(matches)
try:
    precision = float(matches_count) / translation_length
    recall = float(matches_count) / reference_length
    fmean = (precision * recall) / (alpha * precision + (1 - alpha) *
↪recall)
    chunk_count = float(_count_chunks(matches))
    frag_frac = chunk_count / matches_count
except ZeroDivisionError:
    return 0.0
penalty = gamma * frag_frac ** beta
return (1 - penalty) * fmean

def meteor_score(
    references,
    hypothesis,
    preprocess=str.lower,
    stemmer=PorterStemmer(),
    wordnet=wordnet,
    # Original parameter values:
    #alpha=0.9,
    #beta=3,
    #gamma=0.5,
    # Parameter values we chose based on the paper: https://www.cs.cmu.edu/~alavie/METEOR/pdf/Lavie-Agarwal-2007-METEOR.pdf
↪~alavie/METEOR/pdf/Lavie-Agarwal-2007-METEOR.pdf
    alpha=0.81,
    beta=0.83,
    gamma=0.28,
):
    """
    Calculates METEOR score for hypothesis with multiple references as
    described in "Meteor: An Automatic Metric for MT Evaluation with
    HighLevels of Correlation with Human Judgments" by Alon Lavie and
    Abhaya Agarwal, in Proceedings of ACL.
    http://www.cs.cmu.edu/~alavie/METEOR/pdf/Lavie-Agarwal-2007-METEOR.pdf

    In case of multiple references the best score is chosen. This method

```

iterates over `single_meteor_score` and picks the best pair among all the references for a given hypothesis

```
>>> hypothesis1 = 'It is a guide to action which ensures that the military_
↳always obeys the commands of the party'
>>> hypothesis2 = 'It is to insure the troops forever hearing the activity_
↳guidebook that party direct'

>>> reference1 = 'It is a guide to action that ensures that the military_
↳will forever heed Party commands'
>>> reference2 = 'It is the guiding principle which guarantees the military_
↳forces always being under the command of the Party'
>>> reference3 = 'It is the practical guide for the army always to heed the_
↳directions of the party'

>>> round(meteor_score([reference1, reference2, reference3], hypothesis1),4)
0.7398
```

If there is no words match during the alignment the method returns the score as 0. We can safely return a zero instead of raising a division by zero error as no match usually implies a bad translation.

```
>>> round(meteor_score(['this is a cat'], 'non matching hypothesis'),4)
0.0
```

```
:param references: reference sentences
:type references: list(str)
:param hypothesis: a hypothesis sentence
:type hypothesis: str
:param preprocess: preprocessing function (default str.lower)
:type preprocess: method
:param stemmer: nltk.stem.api.StemmerI object (default PorterStemmer())
:type stemmer: nltk.stem.api.StemmerI or any class that implements a stem_
↳method
:param wordnet: a wordnet corpus reader object (default nltk.corpus.wordnet)
:type wordnet: WordNetCorpusReader
:param alpha: parameter for controlling relative weights of precision and_
↳recall.
:type alpha: float
:param beta: parameter for controlling shape of penalty as a function
of as a function of fragmentation.
:type beta: float
:param gamma: relative weight assigned to fragmentation penalty.
:type gamma: float
:return: The sentence-level METEOR score.
:rtype: float
```

```

"""
return max(
    [
        single_meteor_score(
            reference,
            hypothesis,
            stemmer=stemmer,
            wordnet=wordnet,
            alpha=alpha,
            beta=beta,
            gamma=gamma,
        )
        for reference in references
    ]
)

```

[21]: *# Functiones defined for calculating the meteor scores*

```

# Define the function to calculate BLEU scores for more than one inputs
def meteor_lists_no_df(golden_standard, text1, text2):
    # n-gram lists for text1 and text2
    meteor_list = [[], []]

    # Append meteor scores - call the meteor function on text1 and text2
    meteor_list[0].append(meteor_score([golden_standard], text1))
    meteor_list[1].append(meteor_score([golden_standard], text2))

    # This one returns ngram_list directly for the function that makes a row
    → per answer
    return(meteor_list)

def meteor_more(gold1, berts, distilberts):
    # Initiate results dataframe
    results = pd.DataFrame()

    # Loop over rows in the dataframe
    for i in range(0, len(gold1)):

        # Define the inputs to the bleu function from the dataset
        gold_1 = gold1[i]
        bert = berts[i]
        distil = distilberts[i]

        # Get the bleu scores per line in dataframe with the predefined function
        meteor_list1 = meteor_lists_no_df(gold_1, bert, distil)

        # Append to dataframe, index is model name + iteration

```

```

    dff = pd.DataFrame(meteor_list1, index=['0', '1'], columns =
↳ ["METEOR"])

    # Append results to the dataframe
    results = results.append(dff)

    return(results)

```

```
[22]: meteor_more(df['Gold_1'], df['BERT'], df['DistilBERT'])
```

```

[22]:      METEOR
0    0.842492
1    0.842492
0    0.000000
1    0.000000
0    0.000000
..      ...
1    0.000000
0    0.720000
1    0.720000
0    0.328767
1    0.328767

[22130 rows x 1 columns]

```

1.3 2. Calculate GLEU scores

1.3.1 2.1 Prepare the dataframe

```

[58]: import pandas as pd
      # Read in data
      data = pd.read_csv("../tweetQA_bothpresent.csv")

```

```

[59]: # Select columns relevant
      df = data[['Answer', "L_BERT_answer", "DistilBERT_answer"]]
      df = df.rename(columns={'Answer': 'Answers', "L_BERT_answer": "BERT",
↳ "DistilBERT_answer": "DistilBERT"})

```

```
[60]: df
```

```

[60]:      Answers \
0      ['w nj', 'w nj']
1      ['#endangereddriver', '#endangereddriever']
2      ['wiggins', 'wiggins']
3      ['the game is tied at 106', '106-106']
4      ["kemba's", "kemba's floater"]
...      ...

```

```

11060                                ['guns']
11061                                ['president obama']
11062                                ['our best, whole foods']
11063                                ['january']
11064                                ['shed it']

                                BERT \
0                                w nj
1                                # endangeredriver
2    monstars basketball @ m0nstarsbballwiggins
3                                106 - 106
4                                kemba
...                                ...
11060                                guns
11061                                president obama
11062                                it happens to the best of us
11063                                january
11064                                shed their skin

                                DistilBERT
0                                w nj
1                                jdsutter
2    monstars basketball @ m0nstarsbballwiggins
3                                106 - 106 . 8 . 9
4                                kemba
...                                ...
11060                                guns
11061                                president obama
11062                                it happens to the best of us
11063                                january
11064                                shed their skin

[11065 rows x 3 columns]

```

```
[61]: df.describe()
```

```

[61]:
count          11065  BERT      11065  DistilBERT      11065
unique           9432  BERT      8875  DistilBERT      8991
top    ['donald trump']  trump  donald j . trump
freq              33      30              26

```

The answers have sometimes more than one correct option: make the answers into two columns, Gold_1 and Gold_2.

```
[62]: import ast
```

```

# Function to split the rows
def split_column(row):
    # Cuts the answers row into two if possible
    initial_gold = ast.literal_eval(row['Answers'])
    # Gold_1 will always have an input
    row['Gold_1'] = initial_gold[0]
    # If the list has more than 1 element, Gold_2 gets the second input
    if len(initial_gold) > 1:
        row['Gold_2'] = initial_gold[1]

    return(row)

# Apply on the dataframe
df = df.apply(split_column, axis = 1)

```

1.3.2 2.2 Apply the GLEU function

```

[27]: import math
gleus = gleu_more(df['Gold_1'], df['Gold_2'], df['BERT'], df['DistilBERT'])

```

```

[28]: #Define the bert and distilbert results based on the indexes given to them in
      ↪ the function above
bert_results = gleus.loc["0"]
distil_results = gleus.loc["1"]

```

```

[97]: # inspect
bert_results.describe()

```

```

[97]:
count    0
count    11065.000000
mean      0.542009
std       0.436460
min       0.000000
25%       0.055556
50%       0.500000
75%       1.000000
max       1.000000

```

```

[98]: distil_results.describe()

```

```

[98]:
count    0
count    11065.000000
mean      0.472871
std       0.439373
min       0.000000
25%       0.000000
50%       0.333333

```



```
75%      1.000000
max       1.000000
```

```
[345]: bert_results.to_csv("bert_gleus.csv", index = False)
       distil_results.to_csv("distil_gleus.csv", index = False)
```

1.4 3. Calculate METEOR scores

```
[63]: meteors = meteor_more(df['Gold_1'], df['BERT'], df['DistilBERT'])
```

```
[67]: # Get the scores for bert and distilbert
       bert_meteors = meteors.loc["0"]
       distil_meteors = meteors.loc["1"]
```

```
[81]: bert_meteors.describe()
```

```
[81]:          METEOR
count  11065.000000
mean      0.460381
std       0.337903
min       0.000000
25%      0.000000
50%      0.547954
75%      0.720000
max       1.000000
```

The METEOR scores are sometimes above 1. Inspection below in part 4 shows that the score being one usually means the model was correct. Thus, the scores above 1 will be replaced with a score of 1.

```
[89]: bert_meteors['METEOR'] = np.where(bert_meteors['METEOR'] > 1, 1,
    ↪ bert_meteors['METEOR'])
       bert_meteors.describe()
```

```
[89]:          METEOR
count  11065.000000
mean      0.460381
std       0.337903
min       0.000000
25%      0.000000
50%      0.547954
75%      0.720000
max       1.000000
```

```
[90]: distil_meteors['METEOR'] = np.where(distil_meteors['METEOR'] > 1, 1,
    ↪ distil_meteors['METEOR'])
       distil_meteors.describe()
```

```
[90]: METEOR
count 11065.000000
mean   0.415755
std    0.343811
min    0.000000
25%    0.000000
50%    0.397790
75%    0.720000
max    1.000000
```

```
[83]: bert_meteors.to_csv("bert_meteors.csv", index = False)
      distil_meteors.to_csv("distil_meteors.csv", index = False)
```

1.5 4. Appending to dataframe, getting 50 samples

```
[35]: df
```

```
[35]:
      Answers \
0      ['w nj', 'w nj']
1      ['#endangeredriver', '#endangeredriver']
2      ['wiggins', 'wiggins']
3      ['the game is tied at 106', '106-106']
4      ["kemba's", "kemba's floater"]
...
11060      ['guns']
11061      ['president obama']
11062      ['our best, whole foods']
11063      ['january']
11064      ['shed it']

      BERT \
0      w nj
1      # endangeredriver
2      monstars basketball @ m0nstarsbballwiggins
3      106 - 106
4      kemba
...
11060      guns
11061      president obama
11062      it happens to the best of us
11063      january
11064      shed their skin

      DistilBERT      Gold_1 \
0      w nj      w nj
1      jdsutter      #endangeredriver
2      monstars basketball @ m0nstarsbballwiggins      wiggins
```

```

3          106 - 106 . 8 . 9 the game is tied at 106
4          kemba          kemba's
...          ...          ...
11060          guns          guns
11061          president obama          president obama
11062          it happens to the best of us          our best, whole foods
11063          january          january
11064          shed their skin          shed it

```

```

          Gold_2
0          w nj
1          #endangereddriver
2          wiggins
3          106-106
4          kemba's floater
...          ...
11060          NaN
11061          NaN
11062          NaN
11063          NaN
11064          NaN

```

[11065 rows x 5 columns]

```

[84]: # Reset indexes for results
bert_meteors = bert_meteors.reset_index(drop=True)
distil_meteors = distil_meteors.reset_index(drop=True)
bert_results = bert_results.reset_index(drop=True)
distil_results = distil_results.reset_index(drop=True)

# Append to dataframe
df['BERT_METEOR'] = bert_meteors['METEOR']
df['DistilBERT_METEOR'] = distil_meteors['METEOR']
df['BERT_GLEU'] = bert_results[0]
df['DistilBERT_GLEU'] = distil_results[0]

```

```

[85]: df

```

```

[85]:          Answers \
0          ['w nj', 'w nj']
1          ['#endangereddriver', '#endangereddriver']
2          ['wiggins', 'wiggins']
3          ['the game is tied at 106', '106-106']
4          ["kemba's", "kemba's floater"]
...          ...
11060          ['guns']
11061          ['president obama']

```

11062	['our best, whole foods']
11063	['january']
11064	['shed it']

	BERT \
0	w nj
1	# endangeredriver
2	monstars basketball @ m0nstarsbballwiggins
3	106 - 106
4	kemba
...	...
11060	guns
11061	president obama
11062	it happens to the best of us
11063	january
11064	shed their skin

	DistilBERT	Gold_1 \
0	w nj	w nj
1	jdsutter	#endangeredriver
2	monstars basketball @ m0nstarsbballwiggins	wiggins
3	106 - 106 . 8 . 9	the game is tied at 106
4	kemba	kemba's
...
11060	guns	guns
11061	president obama	president obama
11062	it happens to the best of us	our best, whole foods
11063	january	january
11064	shed their skin	shed it

	Gold_2	BERT_METEOR	DistilBERT_METEOR	BERT_GLEU \
0	w nj	0.842492	0.842492	1.000000
1	#endangeredriver	0.000000	0.000000	1.000000
2	wiggins	0.000000	0.000000	0.000000
3	106-106	0.132597	0.116317	0.055556
4	kemba's floater	0.000000	0.000000	0.000000
...
11060	NaN	0.720000	0.720000	1.000000
11061	NaN	0.842492	0.842492	1.000000
11062	NaN	0.000000	0.000000	0.045455
11063	NaN	0.720000	0.720000	1.000000
11064	NaN	0.328767	0.328767	0.166667

	DistilBERT_GLEU
0	1.000000
1	0.000000
2	0.000000

```

3          0.055556
4          0.000000
...
11060      1.000000
11061      1.000000
11062      0.045455
11063      1.000000
11064      0.166667

```

[11065 rows x 9 columns]

```

[102]: # Random sample 50 sentences
samples = df.sample(n=100)

# Read out samples to inspect them in Excel
samples.to_csv("df_samples_scores100.csv")

```

```

[87]: # Make the two dataframes for short and long answers
import numpy as np

short_answers = df.loc[np.array(list(map(len,df["Gold_1"].str.split())) < 3)]
long_answers = df.loc[np.array(list(map(len,df["Gold_1"].str.split())) >= 3)]

```

```

[88]: short_answers.to_csv("df_short_answers.csv")
long_answers.to_csv("df_long_answers.csv")

```

```

[52]: # Inspecting the dataframe where meteor scores were bigger than 1
df[df["BERT_METEOR"] >= 1]

```

```

[52]:
           Answers                                     BERT \
449      ['shark', 'sharks']                                sharks
1027     ['steal from us']                            stealing from us
1099     ['paul walker']                                paul walkers
1418     ['cairnes store']                            cairns stores
1581  ['hundred of thousands of dollars']  hundreds of thousands of dollars
2365     ['smile']                                       smiled
2422  ['transgendered people']                transgender people
3731     ['play racist call']                    plays racist calls
4320     ['stereotype']                          stereotypes
4769     ['work really hard']                    works really hard
5068     ['female action star']                female action stars
5332     ['stop to listen']                    stopping to listen
5765     ['end']                                  ended
7066     ['gun']                                  guns
7373     ['42 degree celsius']                42 degrees celsius
7936     ['help deliver babies']              helps delivers babies
8102     ['black makeup artist']              black makeup artists

```

8362	['dodge bullets']	dodged bullets
8611	['paul walker']	paul walkers
9685	['killing me']	trying to kill me
10666	['masturbate']	masturbated at me
10922	['pull it']	pulling it

	DistilBERT \
449	jason demers (@ jasondemers5) july 16 , 2014
1027	stealing
1099	paul walkers
1418	cairns
1581	hundreds of thousands of dollars
2365	smiled
2422	transgender people
3731	plays racist calls
4320	stereotypes
4769	works really hard
5068	female action stars
5332	stopping to listen
5765	ended
7066	kim kardashian west
7373	42 degrees celsius
7936	helps delivers babies
8102	black makeup artists
8362	dodged bullets
8611	kevin hart
9685	trying to kill me
10666	masturbated
10922	pulling

	Gold_1	Gold_2	BERT_METEOR \
449	shark	sharks	1.440000
1027	steal from us	NaN	1.123322
1099	paul walker	NaN	1.200019
1418	cairnes store	NaN	1.200019
1581	hundred of thousands of dollars	NaN	1.065002
2365	smile	NaN	1.440000
2422	transgendered people	NaN	1.200019
3731	play racist call	NaN	1.361264
4320	stereotype	NaN	1.440000
4769	work really hard	NaN	1.123322
5068	female action star	NaN	1.123322
5332	stop to listen	NaN	1.123322
5765	end	NaN	1.440000
7066	gun	NaN	1.440000
7373	42 degree celsius	NaN	1.123322
7936	help deliver babies	NaN	1.361264

8102	black makeup artist	NaN	1.123322
8362	dodge bullets	NaN	1.200019
8611	paul walker	NaN	1.200019
9685	killing me	NaN	1.008419
10666	masturbate	NaN	1.043478
10922	pull it	NaN	1.200019

	DistilBERT_METEOR	BERT_GLEU	DistilBERT_GLEU
449	0.000000	1.000000	0.000000
1027	0.549618	0.500000	0.000000
1099	1.200019	0.333333	0.333333
1418	0.397790	0.000000	0.000000
1581	1.065002	0.714286	0.714286
2365	1.440000	0.000000	0.000000
2422	1.200019	0.333333	0.333333
3731	1.361264	0.166667	0.166667
4320	1.440000	0.000000	0.000000
4769	1.123322	0.500000	0.500000
5068	1.123322	0.500000	0.500000
5332	1.123322	0.500000	0.500000
5765	1.440000	0.000000	0.000000
7066	0.000000	0.000000	0.000000
7373	1.123322	0.333333	0.333333
7936	1.361264	0.166667	0.166667
8102	1.123322	0.500000	0.500000
8362	1.200019	0.333333	0.333333
8611	0.000000	0.333333	0.000000
9685	1.008419	0.100000	0.100000
10666	1.440000	0.000000	0.000000
10922	0.795580	0.333333	0.000000

```
[99]: short_answers.describe()
```

```
[99]:
```

	BERT_METEOR	DistilBERT_METEOR	BERT_GLEU	DistilBERT_GLEU
count	7245.000000	7245.000000	7245.000000	7245.000000
mean	0.474589	0.430672	0.614320	0.534981
std	0.343093	0.351143	0.447432	0.459766
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.043478	0.000000
50%	0.657534	0.521739	1.000000	0.333333
75%	0.720000	0.720000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

```
[100]: long_answers.describe()
```

```
[100]:
```

	BERT_METEOR	DistilBERT_METEOR	BERT_GLEU	DistilBERT_GLEU
count	3820.000000	3820.000000	3820.000000	3820.000000

mean	0.433435	0.387465	0.404865	0.355074
std	0.326184	0.327642	0.378705	0.370297
min	0.000000	0.000000	0.000000	0.000000
25%	0.146568	0.000000	0.071429	0.000000
50%	0.426683	0.350308	0.300000	0.200000
75%	0.724701	0.665626	0.714286	0.600000
max	1.000000	1.000000	1.000000	1.000000

1.6 Additional code: how GLEU compares to METEOR in short examples

```
[39]: # Meteor is lower for an exact short match
print(meteor_score(["one"], "one"))
print(nltk.translate.gleu_score.sentence_gleu(["one"], "one"))
```

```
0.72
1.0
```

```
[40]: # Meteor is higher for when there are more words in exact match
a = "one and seventeen"
print(meteor_score([a], a))
print(nltk.translate.gleu_score.sentence_gleu([a.split()], a.split()))
```

```
0.8875013295249914
1.0
```